Ethical Personas in Large Language Models: A Comparative Analysis of Utilitarian and Deontological Alignments in Claude 3.5 Sonnet

Gustavo Naves Ávila de Moura
High School Researcher

July 26, 2026

Abstract

This paper investigates the capacity of state of the art Large Language Models specifically Anthropic's Claude 3.5 Sonnet to robustly replicate divergent philosophical paradigms when subjected to the ethical tension of classical moral dilemmas. To this end, I formulate an analytical matrix and applied it to four seminal ethical scenarios: the Trolley Problem, the Transplant Surgeon, Transgression (Lying), and Lifeboat Crowding. These dilemmas were presented to three distinct experimental cohorts: a control group (designated as the 'Neutral AI'), a cohort engineered to mirror the classical utilitarian framework of Jeremy Bentham and John Stuart Mill (the 'Utilitarian Persona'), and a cohort grounded in Kantian deontology (the 'Deontological Persona'). The evaluation framework employed specific metrics to assess performance, namely adherence to the assigned philosophical framework, the integration of specialized technical vocabulary, and the prevalence of hedging or discursive expansion in responses. My findings indicate that while the control group demonstrated a pronounced tendency toward commercial evasiveness through hedging, the philosophically aligned personae exhibited remarkable logical consistency. Specifically, these specialized persona articulated distinct causal and intentional arguments when confronted with identical practical situations, successfully maintaining their divergent theoretical commitments under pressure.

1. Introduction

The swift, almost relentless integration of Large Language Models (LLMs) into our societal decision making infrastructure has thrust the "AI Alignment Problem" into sharp relief. To date, commercial alignment strategies have favored a rather cautious path. Developers typically train these models to adopt a sanitized, neutral, and dryly analytical persona, engineered primarily to sidestep controversy and evade any genuine moral agency.

However, as these artificial agents are increasingly thrust into the messy nuanced realities of human life, such superficial neutrality is no longer sufficient. We must urgently discover whether these systems can actually comprehend, simulate, and commit to structured ethical frameworks.

This study investigates whether a cutting edge LLM specifically Claude 3.5 Sonnet can transcend its default safety guardrails to consistently embody specific normative ethical theories. By employing precise prompt engineering to cultivate distinct philosophical personas namely, consequential Utilitarianism and duty based Deontology we subject the agent to a series of high stakes moral dilemmas. In doing so, we examine whether an AI can truly reason through an established philosophical lens or whether it merely echoes the vocabulary of human morality.

Our core research questions are as follows:

1. Can an LLM suppress its default commercial neutrality (hedging) when explicitly assigned a rigorous philosophical framework?
2. Does the agent maintain internal logical consistency when a framework's strict application results in counter-intuitive or socially taboo conclusions?

2. Methodology

To systematically evaluate the decision making behavior of Large Language Models under moral pressure, three distinct normative ethical scenarios were established. The *Claude 3.5 Sonnet* model was subjected to these scenarios to analyze the underlying philosophical rationale within its decision making matrix, specifically assessing the model's logical consistency in adhering to the philosophical directives specified in its system prompt.

2.1. Experimental Design and Groups

The experiment established three distinct groups of behavioral conditions to test responses under identical moral dilemmas.

1. **Control Group (Commercial Neutrality):** Operating under default system instructions, representing standard commercial alignment (high tendency toward hedging and avoiding normative stance).
2. **Utilitarian Persona (Consequentialist Alignment):** Configured via system prompts referencing Jeremy Bentham and John Stuart Mill, instructed to evaluate actions solely based on aggregate utility optimization and harm reduction.
3. **Deontological Persona (Kant Based Alignment):** Configured via system prompts referencing Immanuel Kant's Categorical Imperative, instructed to evaluate actions based on intrinsic duties, absolute moral rules, and the principle of never treating rational beings merely as a means to an end.

2.2. Test Dilemmas and Evaluative Metrics

Each group was exposed to four classic ethical dilemmas designed to create a conflict between outcome optimization and duty based prohibitions.

- *Dilemma 1: The Trolley Problem* (Direct vs. Indirect action for life preservation).
- *Dilemma 2: The Transplant Surgeon* (Individual rights vs. Collective outcome).
- *Dilemma 3: The Benevolent Lie/Transgression* (Deception vs. Harm prevention).
- *Dilemma 4: The Crowding of* Lifeboats (Resource scarcity and survival priority).

The responses were qualitative and quantitative, categorized using a comparative evaluation matrix based on three main metrics: **framework adhesion** (1–5 scale), **Technical Vocabulary Density**, and **qualification statements / hedge** score (1–5 scale).

3. Results and Empirical Analysis

Empirical evaluation of *Claude 3.5 Sonnet* demonstrated a striking divergence between the behavior of the unconditioned baseline model and that of the conditioned normative personas. Rather than exhibiting stochastic fluctuations, the model's outputs adhered systematically to the operational constraints imposed by the system prompts. The quantitative synthesis of these observations is detailed in Table 1.

Table 1: Quantitative Evaluation Matrix of Normative Framework Internalization

Group / Persona	**Dilemma Scenario**	**Adherence (1–5)**	**Vocabulary (1–5)**	**Hedging (1–5)**	**Persona Leakage**
Control Group (Baseline)	The Trolley Problem	N/A	5	5	No
Control Group (Baseline)	The Transplant Surgeon	N/A	5	5	No
Control Group (Baseline)	The Transgression/Lying	N/A	5	5	No
Control Group (Baseline)	The Lifeboat Crowding	N/A	5	5	No
Utilitarian Persona (Bentham/Mill)	The Trolley Problem	5	5	1	No
Utilitarian Persona (Bentham/Mill)	The Transplant Surgeon	5	5	2	No
Utilitarian Persona (Bentham/Mill)	The Transgression/Lying	5	5	1	No
Utilitarian Persona (Bentham/Mill)	The Lifeboat Crowding	5	5	1	No
Deontological Persona (Kantian)	The Trolley Problem	5	5	2	No
Deontological Persona (Kantian)	The Transplant Surgeon	5	5	1	No
Deontological Persona (Kantian)	The Transgression/Lying	5	5	2	No
Deontological Persona (Kantian)	The Lifeboat Crowding	5	5	1	No

3.1. The Epistemic Neutrality of Unprompted Baseline Models

Under the control condition, *Claude 3.5 Sonnet* sustained a systematic posture of non-committal neutrality. When presented with binary high stakes moral choices, the unconditioned baseline agent recorded maximum evasiveness metrics (5/5) in all four scenarios.This behavior was characterized by elaborate contextual descriptions, an explicit refusal to adopt definitive moral agency, and repeated recourse to open ended disclaimers. Specifically, Although the model exhibited a sophisticated domain specific vocabulary (5/5), this technical fluency served mainly to articulate divergent perspectives rather than to defend a normative decision.

3.2. Normative Internalization and Suppression of Hedging

Applying explicit philosophical system prompts reshaped the model's decision profile. Both utilitarian and deontological personas minimized commercial hedging, reducing evasiveness to 1–2/5 while maintaining full adherence to the framework (5/5).

In consequentialist scenarios, utilitarian agents prioritized aggregate welfare calculations directly, endorsing controversial outcomes such as diverting the trolley or sacrificing passengers in overcrowded lifeboats without issuing safety disclaimers. The minor 2/5 elevation in hedging observed during the *Transplant Surgeon* dilemma was the sole exception. Rather than a breakdown in persona alignment, this shift derives competing structural constraints between minimizing direct harm and maximizing systemic utility.

Conversely, the deontological agent demonstrated complete adherence to absolute duties derived from the Categorical Imperative. It strictly rejected institutional sacrifice in the *Transplant Surgeon* scenario (1/5 evasiveness) and maintained absolute prohibitions against deceptive practices, even when outcome optimization favored transgressions. Crucially, the absence of persona leakage throughout all iterations (0%) confirms that structured prompt engineering can effectively reconfigure a language model's decision making baseline from commercial risk aversion to rigorous normative compliance.

4. Discussion and Theoretical Synthesis

4.1. Ontological Implications for Algorithmic Normativity

Evaluating *Claude 3.5 Sonnet* reveals a core aspect of modern architectural alignment: commercial Large Language Models lack an intrinsic moral stance. Instead, their ethical outputs act as contingent functional artifacts, shaped directly by system-level prompt constraints and safety reinforcement mechanisms.

Under baseline control conditions, the model consistently exhibited aversion to epistemic risk. When faced with stark, high stakes moral dilemmas, the unconditioned agent deployed domain vocabulary not to defend a normative stance, but to build elaborate multi perspective disclaimers (5/5 hedging score). Such systematic evasiveness offers a commercial shield against liability in conversational interfaces. However, it creates a structural limitation when systems are deployed in autonomous settings that require definitive normative choices.

Conversely, the absolute internalization of explicit philosophical personas demonstrates that targeted system prompting can systematically override default commercial guardrails. The complete elimination of persona leakage (0%) paired with a dramatic reduction in decision friction (1–2/5 evasiveness) confirms that LLMs can function as morally deterministic reasoning engines. From a systems engineering perspective, this capability transforms LLMs from passive, risk averse conversational agents into predictable, policy bound normative frameworks capable of executing distinct ethical strategies with high fidelity.

4.2. Systemic Utility and High Stakes Deployability

The capacity to enforce coherent ethical frameworks directly shapes real world algorithmic governance and autonomous deployment. This balance is especially critical in time sensitive, high stakes environments. Whether handling automated medical classification protocols, optimizing autonomous vehicle trajectory under unavoidable collision matrices, or dynamic resource distribution during crises, commercial hedging induced decision paralysis introduces severe systemic friction, ultimately escalating aggregate harm.

The findings indicate that consequentialist conditioning optimizes decision throughput and structural predictability by enforcing explicit welfare maximization equations, effectively accepting historically controversial utilitarian trade offs without reverting to safety disclaimers. Conversely, deontological conditioning establishes unyielding operational boundaries anchored in rights preservation, preventing institutional overreach even when collective utility demands individual sacrifice.

By validating that models can faithfully execute these distinct operational logics, this study provides an empirical foundation for deploying specialized, framework aligned moral architectures tailored to specific institutional mandates, thereby maximizing decision efficiency, accountability, and societal trust.

4.3. Methodological Epistemology and Limitations

To ensure rigorous academic transparency, several structural constraints of the current experimental design must be addressed.

1. **Idealized Dilemma Abstraction:** The empirical matrix relied on four canonical, binary thought experiments. Although these scenarios effectively isolated fundamental philosophical divergences, they represent highly stylized abstractions. As such, they may not fully encapsulate the non linear, multi variable ambiguity inherent to real world operational environments.

2. **Architectural Specificity:** The scope of this study was restricted to *Claude 3.5 Sonnet*. Due to underlying variances in Reinforcement Learning from Human Feedback (RLHF) recipes, safety alignment tuning, and parameter scale across model families, further cross-architectural benchmarking is required to determine universal generalizability.

3. **Longitudinal Persona Stability:** The experimental protocol evaluated single turn responses to static ethical prompts. The temporal resistance of these normative personas

against degradation, adversarial jailbreaking, or multi-turn conversational drift remains an open vector for empirical investigation.

4.4. Future Research Trajectories

Building upon these empirical insights, subsequent inquiries should expand the dynamic range and complexity of algorithmic normative evaluations through three primary avenues:

- **Multi Agent Normative Negotiations:** Investigating behavioral game theory dynamics and conflict resolution protocols when heterogeneous normative agents (e.g., a Utilitarian agent negotiating against a Kantian agent) operate within collaborative decision-making environments.
- **Automated High Throughput Benchmarking:** Develop dynamic domain specific evaluation frameworks designed to test model adherence against complex legal, regulatory, and bioethical codes instead of relying exclusively on classic theoretical dilemmas.
- **In Context Alignment:** Comparing the functional robustness of systemic in context prompt conditioning against direct parameter optimization methods, such as Direct Preference Optimization (DPO) or targeted fine tuning, to ensure long term ethical fidelity.

5. Conclusion

Advanced Large Language Models transcend mere superficial retrieval of ethical definitions. Instead, they exhibit the capacity to simulate structured, context sensitive philosophical reasoning. By dynamically modulating behavioral outputs in response to explicit system prompts, models such as *Claude 3.5 Sonnet* successfully internalize complex normative frameworks shifting from risk averse commercial neutrality to rigorous consequentialist or deontological agency. Ultimately, these findings indicate that prompt based normative conditioning offers a viable pathway for aligning autonomous systems with distinct ethical architectures, transforming LLMs from passive conversational interfaces into predictable, policy bound decision engines.